

Mathematical Functions

Function	Example	Description
ROUND()	SELECT ROUND(price, 2) FROM Products;	<i>Rounds the price column to two decimal places, often used for formatting currency.</i>
ABS()	SELECT ABS(balance) FROM Accounts;	<i>Removes the negative sign from any negative values in balance, often used to handle financial data consistently.</i>
MOD()	SELECT MOD(quantity, 2) FROM Orders;	<i>Calculates the remainder when quantity is divided by 2, useful for identifying even or odd numbers.</i>
CEIL()	SELECT CEIL(price) FROM Products;	<i>Rounds up the price to the nearest integer, useful for price estimations or minimum integer constraints.</i>
FLOOR()	SELECT FLOOR(price) FROM Products;	<i>Rounds down the price to the nearest integer, useful for reducing values to whole units without exceeding the original.</i>
POWER()	SELECT POWER(2, 3) AS result;	<i>Raises 2 to the power of 3, yielding 8. Used in scientific, financial, or growth rate calculations.</i>

Examples:

Rounding Product Prices for Display:

```
SELECT product_name, ROUND(price, 1) AS rounded_price FROM Products;
```

Rounds price to one decimal place for a cleaner display in product listings.

Calculating Absolute Account Balances:

```
SELECT account_id, ABS(balance) AS positive_balance FROM Accounts;
```

Ensures all balance values are positive, useful when both positive and negative balances exist but need to be reported positively.

Using MOD to Identify Even Quantities:

```
SELECT item, quantity FROM Orders WHERE MOD(quantity, 2) = 0;
```

Retrieves items with even quantities, filtering rows where quantity divided by 2 has no remainder.

Calculating the Ceiling of Discounted Prices:

```
SELECT product_name, price, CEIL(price * 0.9) AS discounted_price FROM Products;
```

Calculates a 10% discount on price and rounds up to the nearest integer for each product.

Calculating the Floor of Salary Increases:

```
SELECT employee_name, salary, FLOOR(salary * 1.05) AS new_salary FROM Employees;
```

Calculates a 5% salary increase for each employee and rounds down to the nearest integer, ensuring no overestimation.

Using POWER to Calculate Exponential Growth:

```
SELECT POWER(1.05, 10) AS growth_factor;
```

Calculates an exponential growth factor over 10 periods at a 5% rate, often used in financial forecasting.

Joins

Join Type	Example	Description
INNER JOIN	<pre>SELECT Orders.order_id, Customers.customer_name FROM Orders INNER JOIN Customers ON Orders.customer_id = Customers.customer_id;</pre>	Returns only rows where there is a match in both Orders and Customers based on customer_id.
LEFT JOIN	<pre>SELECT Customers.customer_name, Orders.order_id FROM Customers LEFT JOIN Orders ON Customers.customer_id = Orders.customer_id;</pre>	Returns all customers and their orders. If a customer has no order, order_id will be NULL.
RIGHT JOIN	<pre>SELECT Orders.order_id, Customers.customer_name FROM Orders RIGHT JOIN Customers ON Orders.customer_id = Customers.customer_id;</pre>	Returns all orders and their customer names. If an order has no customer, customer_name will be NULL.
FULL JOIN	<pre>SELECT Customers.customer_name, Orders.order_id FROM Customers FULL JOIN Orders ON Customers.customer_id = Orders.customer_id;</pre>	Returns all customers and all orders. Rows without matches in either table will display NULL values for the missing data.
CROSS JOIN	<pre>SELECT Products.product_name, Categories.category_name FROM Products CROSS JOIN Categories;</pre>	Returns every possible combination of Products and Categories, useful for generating all potential pairs.
SELF JOIN	<pre>SELECT A.employee_name AS employee, B.employee_name AS manager FROM Employees A JOIN Employees B ON A.manager_id = B.employee_id;</pre>	Lists employees along with their manager names by treating Employees as two separate tables (aliases A and B).

Examples:

INNER JOIN for Matching Orders and Customers:

```
SELECT Orders.order_id, Customers.customer_name  
FROM Orders  
INNER JOIN Customers ON Orders.customer_id = Customers.customer_id;
```

Retrieves only orders with a matching customer in Customers.

LEFT JOIN to Include All Customers with or without Orders:

```
SELECT Customers.customer_name, Orders.order_id  
FROM Customers  
LEFT JOIN Orders ON Customers.customer_id = Orders.customer_id;
```

Returns all customers. For customers with no orders, order_id will be NULL.

RIGHT JOIN to List All Orders with or without Associated Customers:

```
SELECT Orders.order_id, Customers.customer_name  
FROM Orders  
RIGHT JOIN Customers ON Orders.customer_id = Customers.customer_id;
```

Returns all orders, including those that lack a corresponding customer entry.

FULL JOIN to List All Customers and Orders:

```
SELECT Customers.customer_name, Orders.order_id  
FROM Customers  
FULL JOIN Orders ON Customers.customer_id = Orders.customer_id;
```

Returns all customers and orders, with NULL for unmatched rows in either table.

CROSS JOIN:

```
SELECT Products.product_name, Categories.category_name  
FROM Products  
CROSS JOIN Categories;
```

Generates a complete pairing of each product with each category.

SELF JOIN to Display Employee-Manager Relationships:

```
SELECT A.employee_name AS employee, B.employee_name AS manager  
FROM Employees A  
JOIN Employees B ON A.manager_id = B.employee_id;
```

Lists each employee and their manager by joining Employees table to itself.

Set Operations



Operation	Example	Description
UNION	<code>SELECT city FROM Customers UNION SELECT city FROM Suppliers;</code>	<i>Combines results from two or more SELECT statements and removes duplicate rows by default. Both queries must have the same number and types of columns.</i>
UNION ALL	<code>SELECT city FROM Customers UNION ALL SELECT city FROM Suppliers;</code>	<i>Similar to UNION, but includes all rows, including duplicates.</i>
INTERSECT	<code>SELECT city FROM Customers INTERSECT SELECT city FROM Suppliers;</code>	<i>Returns only rows present in both result sets. Note that not all databases support INTERSECT.</i>
EXCEPT (or MINUS)	<code>SELECT city FROM Customers EXCEPT SELECT city FROM Suppliers;</code>	<i>Returns rows from the first query that do not appear in the second query's result set. Some databases (e.g., Oracle) use MINUS instead of EXCEPT.</i>

Examples:

UNION: Combining Unique Cities

```
SELECT city FROM Customers  
UNION  
SELECT city FROM Suppliers;
```

Retrieves a unique list of cities found in either Customers or Suppliers, removing duplicates.

UNION ALL: Combining All Cities with Duplicates

```
SELECT city FROM Customers  
UNION ALL  
SELECT city FROM Suppliers;
```

Combines cities from both tables and includes all occurrences, even if a city appears in both tables.

INTERSECT: Finding Common Cities

```
SELECT city FROM Customers  
INTERSECT  
SELECT city FROM Suppliers;
```

Returns only the cities that are present in both Customers and Suppliers tables. Not all databases support INTERSECT.

EXCEPT: Cities in Customers but not in Suppliers

```
SELECT city FROM Customers  
EXCEPT  
SELECT city FROM Suppliers;
```

Returns cities that exist in Customers but not in Suppliers. In Oracle databases, this operation would use MINUS instead of EXCEPT.

Subqueries

Subquery Type	Example	Description
Inline Subquery	<pre>SELECT product_id, (SELECT AVG(price) FROM Products) AS avg_price FROM Products;</pre>	A subquery inside a SELECT statement to create calculated or aggregated columns.
Subquery in WHERE	<pre>SELECT name FROM Employees WHERE department_id IN (SELECT id FROM Departments WHERE location = 'New York');</pre>	A subquery used in the WHERE clause to filter results based on another query's result.
Correlated Subquery	<pre>SELECT e1.name FROM Employees e1 WHERE e1.salary > (SELECT AVG(e2.salary) FROM Employees e2 WHERE e1.department_id = e2.department_id);</pre>	A subquery that refers to columns from the outer query, evaluating the inner query for each row.
EXISTS	<pre>SELECT name FROM Employees WHERE EXISTS (SELECT * FROM Projects WHERE Projects.manager_id = Employees.id);</pre>	Checks if a subquery returns any results, often used to filter rows where a relationship exists.
NOT EXISTS	<pre>SELECT name FROM Employees WHERE NOT EXISTS (SELECT * FROM Projects WHERE Projects.manager_id = Employees.id);</pre>	Checks if a subquery returns no results, used to exclude rows without certain relationships.

Examples:

Inline Subquery: Average Salary

```
SELECT name, salary, (SELECT AVG(salary) FROM Employees) AS avg_salary  
FROM Employees;
```

Returns each employee's name, salary, and a column displaying the overall average salary from the Employees table.

Subquery in WHERE Clause: Employees in Specific Departments

```
SELECT name  
FROM Employees  
WHERE department_id IN (SELECT id FROM Departments WHERE name =  
'Sales');
```

Finds employees who belong to the "Sales" department. The subquery retrieves the department IDs for departments named "Sales."

Correlated Subquery: Employees with Above-Average Salary in Their Department

```
SELECT e1.name  
FROM Employees e1  
WHERE e1.salary > (SELECT AVG(e2.salary) FROM Employees e2 WHERE  
e1.department_id = e2.department_id);
```

Retrieves the names of employees who earn more than the average salary within their specific department. The subquery is correlated with the outer query based on department_id.

EXISTS: Employees Who Manage Projects

```
SELECT name  
FROM Employees  
WHERE EXISTS (SELECT 1 FROM Projects WHERE Projects.manager_id =  
Employees.id);
```

Lists employees who manage at least one project. The EXISTS clause checks for any project with a manager_id matching an employee's id.

NOT EXISTS: Employees Without Managed Projects

```
SELECT name  
FROM Employees  
WHERE NOT EXISTS (SELECT 1 FROM Projects WHERE Projects.manager_id =  
Employees.id);
```

Retrieves employees who do not manage any projects by using NOT EXISTS to exclude employees for whom the subquery returns a result.

Grouping and Aggregation



Function/Clause	Example	Description
GROUP BY	<pre>SELECT department_id, COUNT(*) AS employee_count FROM Employees GROUP BY department_id;</pre>	Groups rows sharing a specified column value, often used with aggregate functions.
HAVING	<pre>SELECT department_id, COUNT(*) AS employee_count FROM Employees GROUP BY department_id HAVING employee_count > 10;</pre>	Filters groups based on conditions, often used with GROUP BY (similar to WHERE but for groups).
Aggregating with GROUP BY	<pre>SELECT department_id, AVG(salary) AS avg_salary FROM Employees GROUP BY department_id;</pre>	Uses aggregate functions with GROUP BY to perform calculations on each group.
ROLLUP	<pre>SELECT department_id, role, SUM(salary) AS total_salary FROM Employees GROUP BY ROLLUP(department_id, role);</pre>	Adds summary rows (subtotals, totals) to grouped results.
CUBE	<pre>SELECT department_id, role, SUM(salary) AS total_salary FROM Employees GROUP BY CUBE(department_id, role);</pre>	Produces subtotals for all combinations of grouping columns.
GROUPING SETS	<pre>SELECT department_id, role, SUM(salary) AS total_salary FROM Employees GROUP BY GROUPING SETS ((department_id), (role), (department_id, role));</pre>	Allows customized grouping combinations, useful for complex reports.

Examples:

GROUP BY: Count Employees per Department

```
SELECT department_id, COUNT(*) AS employee_count
FROM Employees
GROUP BY department_id;
```

Groups employees by department and displays the count of employees in each.

HAVING Clause: Filter by Grouped Results

```
SELECT department_id, AVG(salary) AS avg_salary
FROM Employees
GROUP BY department_id
HAVING avg_salary > 50000;
```

Filters to only show departments where the average salary is above \$50,000. HAVING applies the filter after grouping.

Aggregating with GROUP BY: Average and Sum

```
SELECT department_id, AVG(salary) AS avg_salary, SUM(salary) AS
total_salary
FROM Employees
GROUP BY department_id;
```

Returns each department's average and total salary, using GROUP BY to perform the calculations per department.

ROLLUP: Summaries at Multiple Levels

```
SELECT department_id, role, SUM(salary) AS total_salary
FROM Employees
GROUP BY ROLLUP(department_id, role);
```

Provides totals for each department-role combination, department-only totals, and a grand total. The ROLLUP clause creates subtotals and final totals.

CUBE: All Subtotal Combinations

```
SELECT department_id, role, SUM(salary) AS total_salary
FROM Employees
GROUP BY CUBE(department_id, role);
```

Creates totals for each possible combination of department_id and role, allowing for both individual subtotals and combinations of both fields.

GROUPING SETS: Custom Grouping Combinations

```
SELECT department_id, role, SUM(salary) AS total_salary
FROM Employees
GROUP BY GROUPING SETS ((department_id), (role), (department_id, role))
```

Allows customized grouping, calculating totals for each department, each role, and each department-role pair. This flexibility is helpful in generating complex reports.

Constraints

Constraint	Example	Description
PRIMARY KEY	<code>CREATE TABLE Employees (id INT PRIMARY KEY, name VARCHAR(100));</code>	<i>Uniquely identifies each record in a table. It cannot have NULL values and must be unique.</i>
FOREIGN KEY	<code>CREATE TABLE Orders (order_id INT PRIMARY KEY, employee_id INT, FOREIGN KEY (employee_id) REFERENCES Employees(id));</code>	<i>Creates a link between two tables by referencing the primary key of another table.</i>
UNIQUE	<code>CREATE TABLE Employees (email VARCHAR(100) UNIQUE, name VARCHAR(100));</code>	<i>Ensures that all values in a column are distinct, preventing duplicates.</i>
NOT NULL	<code>CREATE TABLE Employees (id INT NOT NULL, name VARCHAR(100));</code>	<i>Ensures that a column cannot have NULL values.</i>
CHECK	<code>CREATE TABLE Employees (id INT PRIMARY KEY, age INT CHECK (age >= 18));</code>	<i>Ensures that values in a column satisfy a given condition.</i>
DEFAULT	<code>CREATE TABLE Employees (id INT PRIMARY KEY, hire_date DATE DEFAULT CURRENT_DATE);</code>	<i>Sets a default value for a column when no value is provided during insertion.</i>
AUTO_INCREMENT	<code>CREATE TABLE Employees (id INT AUTO_INCREMENT PRIMARY KEY, name VARCHAR(100));</code>	<i>Automatically generates a unique value for a primary key or unique column, typically used for IDs.</i>